

# Retrieval Augmented Generation (RAG)

# Agenda

- Retrieval Augmented Augmentation (RAG) Quiz
- Building Blocks of RAG
- Chunking
- Semantic Search
- Evaluation

**Let's begin the discussion by answering a few questions.**

# RAG Quiz

**Which of the following represents the correct sequence of blocks for processing data in a retrieval-augmented generation (RAG) model?**

**A**

Embedding -> Vector DB -> Data Chunk -> Retriever -> LLM

**B**

Data Chunk -> Embedding -> Vector DB -> Retriever -> LLM

**C**

Vector DB -> Data Chunk -> Retriever -> Embedding -> LLM

**D**

Data Chunk -> Embedding -> Retriever -> Vector DB

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# RAG Quiz

**Which of the following represents the correct sequence of blocks for processing data in a retrieval-augmented generation (RAG) model?**

A

Embedding -> Vector DB -> Data Chunk -> Retriever -> LLM

B

Data Chunk -> Embedding -> Vector DB -> Retriever -> LLM

C

Vector DB -> Data Chunk -> Retriever -> Embedding -> LLM

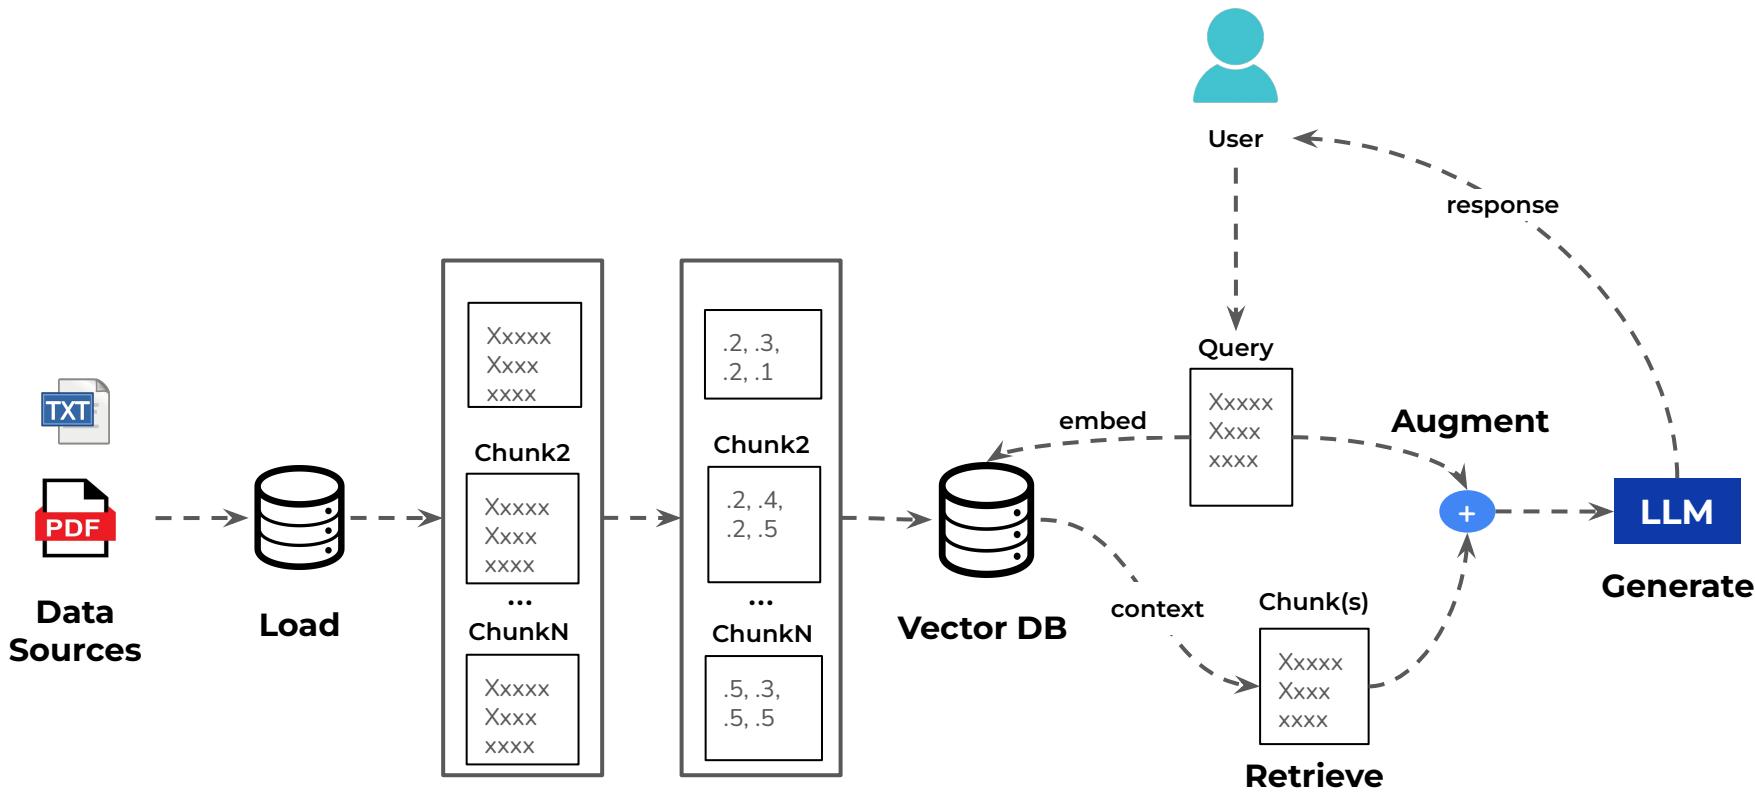
D

Data Chunk -> Embedding -> Retriever -> Vector DB

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Retrieval Augmented Generation (RAG)



This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
 Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# RAG Quiz

## How does increasing the chunk overlap in a RAG system impact retrieval?

A

It reduces the number of retrieved chunks, improving efficiency

B

It increases redundancy, but helps maintain continuity in retrieved content

C

It eliminates the need for chunking altogether

D

It prevents the model from retrieving any irrelevant information

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# RAG Quiz

## How does increasing the chunk overlap in a RAG system impact retrieval?

A

It reduces the number of retrieved chunks, improving efficiency

B

It increases redundancy, but helps maintain continuity in retrieved content

C

It eliminates the need for chunking altogether

D

It prevents the model from retrieving any irrelevant information

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Chunking

The process of **breaking a document** (or text corpus) **into smaller, manageable pieces** (chunks) before storing them

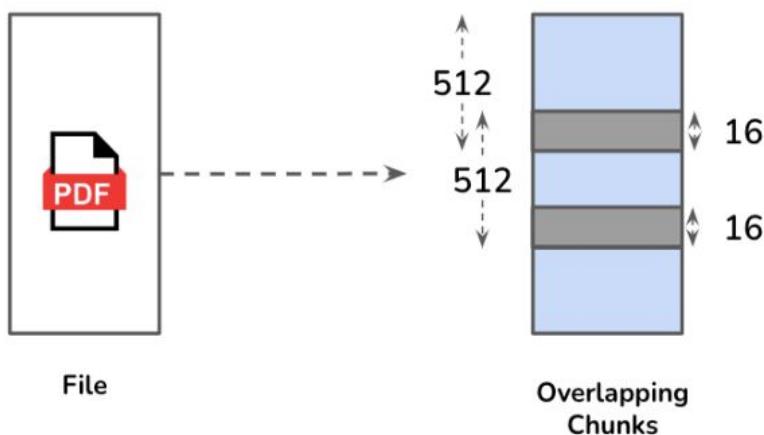
Chunks are then **indexed and used to retrieve relevant information efficiently** when generating responses

512 => number of tokens in a chunk

16 => number of tokens overlapping between any two consecutive chunks

Increasing the chunk overlap parameter increases redundancy as there will be larger overlaps

But this helps maintain continuity in retrieved content as information might be spread across multiple chunks



# RAG Quiz

**Why do we need to create embeddings for our data in a Retrieval-Augmented Generation (RAG) system?**

**A**

To improve keyword-based searching

**B**

To enable semantic similarity search

**C**

To reduce the amount of stored data

**D**

To make data more structured

# RAG Quiz

**Why do we need to create embeddings for our data in a Retrieval-Augmented Generation (RAG) system?**

**A**

To improve keyword-based searching

**B**

To enable semantic similarity search

**C**

To reduce the amount of stored data

**D**

To make data more structured

# Sentence Embeddings

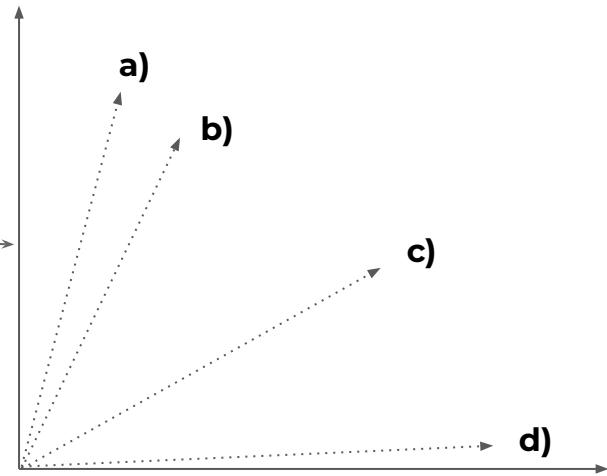
a) Reading books at night is relaxing.

b) Enjoying a good novel before bed is a great habit.

c) The cat jumped over the fence.

d) The weather is cold today.

**Embedding Model**



**Similar sentences are closer in the embedding space; dissimilar sentences are further apart**

This enables **semantic search** to retrieve sentences (or chunks) from the data source that are “similar” to the user query

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# RAG Quiz

**How does RAG enhance the capabilities of LLMs in natural language processing tasks ?**

**A**

By limiting the access to external knowledge sources for generating responses

**B**

By reducing the need for fine-tuning LLMs on specific tasks

**C**

By integrating external knowledge sources to provide contextually rich and accurate responses

**D**

By reducing the model's size

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# RAG Quiz

**How does RAG enhance the capabilities of LLMs in natural language processing tasks ?**

A

By limiting the access to external knowledge sources for generating responses

B

By reducing the need for fine-tuning LLMs on specific tasks

C

By integrating external knowledge sources to provide contextually rich and accurate responses

D

By reducing the model's size

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Advantage of RAG

How many cores does NVIDIA GeForce RTX 5070 Ti have?

**LLM**

I don't have any information on that, as my training data goes up until 2023.

How many cores does NVIDIA GeForce RTX 5070 Ti have?

**RAG**

It has 8960 CUDA cores and a boost clock of 2300 MHz



**NVIDIA GeForce RTX  
Graphics Cards  
Specification  
Document**

# RAG Quiz

**In a Retrieval-Augmented Generation (RAG) system, which factor primarily ensures that the response correctly utilizes the provided information?**

**A**

Clarity in conveying the task

**B**

Relevance of the response to the query

**C**

Faithfulness to the context

**D**

Length of the response

# RAG Quiz

**In a Retrieval-Augmented Generation (RAG) system, which factor primarily ensures that the response correctly utilizes the provided information?**

**A**

Clarity in conveying the task

**B**

Relevance of the response to the query

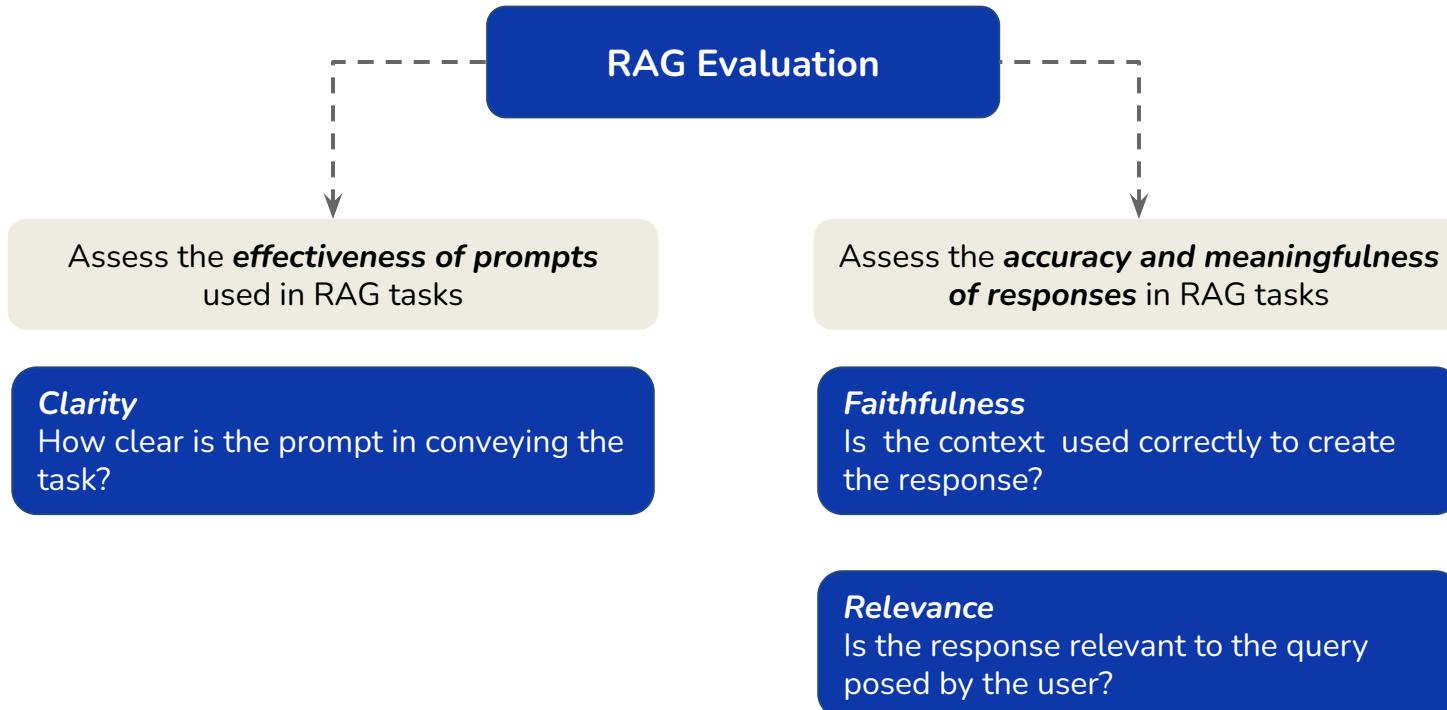
**C**

Faithfulness to the context

**D**

Length of the response

# Evaluation





# Power Ahead!



This file is meant for personal use by erikvdesigner@gmail.com only.  
Sharing or publishing the contents in part or full is liable for legal action.